

# Differentiable K-Means (DKM)

## Soft Clustering for Neural Network Compression

Based on Cho, Vahid, Adya and Rastegari, Apple  
ICLR 2022, arXiv:2108.12659v4

November 6, 2025

# Outline

- 1 Introduction and Motivation
- 2 Differentiable Formulation
- 3 Integration in Neural Networks
- 4 Experimental Results
- 5 Conclusions

# Motivation

- Modern deep networks contain millions of parameters
- Model compression reduces storage, bandwidth, and latency
- Clustering-based compression replaces weights with shared centroid
- Challenge: traditional K-Means is non-differentiable
- Approach taken:
  - original loss function and architecture fixed
  - introduce new layer with no learnable parameters
  - based on attention mechanism to capture cluster interactions
  - weight-based clustering
  - efficient multi-dimensional k-means clustering
  - more than 10x compression factor (parameters)
  - appropriate for both computer vision and language modelling tasks

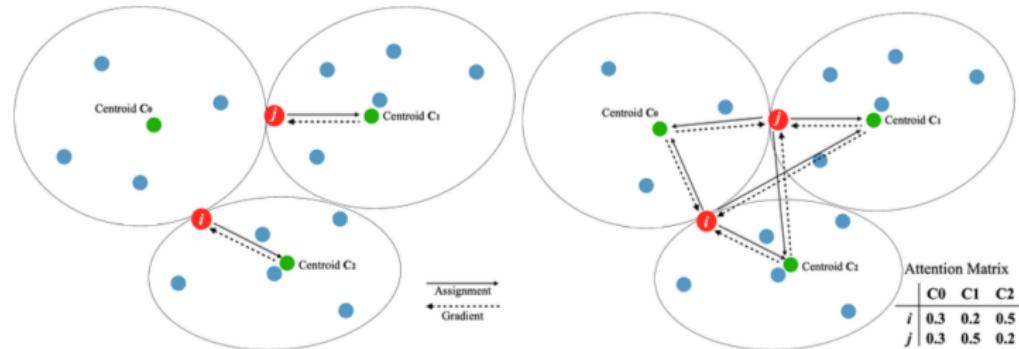
# Motivation

- Modern deep networks contain millions of parameters
- Model compression reduces storage, bandwidth, and latency
- Clustering-based compression replaces weights with shared centroid
- Challenge: traditional K-Means is non-differentiable
- Approach taken:
  - original loss function and architecture fixed
  - introduce new layer with no learnable parameters
  - based on attention mechanism to capture cluster interactions
  - weight-based clustering
  - efficient multi-dimensional k-means clustering
  - more than 10x compression factor (parameters)
  - appropriate for both computer vision and language modelling tasks

## Goal of DKM

Make the clustering process differentiable and trainable end-to-end with the neural network.

# Overcoming limitations



- Conventional approaches make hard assignments for  $i$  and  $j$  based on distance metric
- Gradient is computed only based on the assigned centroid
- Opportunity lost especially with small number of centroids
- Assigning  $i$  to  $C_0$  and  $j$  to  $C_2$  could be better for training loss given small distance difference
- Centroid weight = distance-based attention optimization
- Gradient of a centroid becomes a product of attentions

# Classical K-Means Recap

- $N$  points,  $K$  clusters
- Objective - minimize **inertia**:

$$L_{\text{k-means}} = \sum_{i=1}^N \min_k \|x_i - \mu_k\|^2$$

- Update rules (Lloyd's algorithm)
- Cluster centroids computed based on the  $r_{ik}$  indicator function

$$\mu_k = \frac{\sum_i r_{ik} x_i}{\sum_i r_{ik}}$$

- Non-differentiable because of indicator function:

$$r_{ik} = \begin{cases} 1 & \text{if } k = \arg \min_j \|x_i - \mu_j\|^2 \\ 0 & \text{otherwise} \end{cases}$$

# Soft Assignments in DKM

- Replace hard assignments by soft, differentiable probabilities
- $a_{ik}$  is the **attention** for sample  $i$  w.r.t. centroid  $k$

$$a_{ik} = \frac{\exp(-\frac{\|x_i - \mu_k\|^2}{\tau})}{\sum_j \exp(-\frac{\|x_i - \mu_j\|^2}{\tau})}$$

where  $\alpha = \frac{1}{\tau}$  is the **sharpness (inverse temperature)**.

- Reconstructed weight:

$$\hat{x}_i = \sum_{k=1}^K a_{ik} \mu_k$$

# Soft Assignments in DKM

- Replace hard assignments by soft, differentiable probabilities
- $a_{ik}$  is the **attention** for sample  $i$  w.r.t. centroid  $k$

$$a_{ik} = \frac{\exp(-\frac{\|x_i - \mu_k\|^2}{\tau})}{\sum_j \exp(-\frac{\|x_i - \mu_j\|^2}{\tau})}$$

where  $\alpha = \frac{1}{\tau}$  is the **sharpness (inverse temperature)**.

- Reconstructed weight:

$$\hat{x}_i = \sum_{k=1}^K a_{ik} \mu_k$$

## Temperature Role

Define  $T = 1/\alpha$ .

High  $T \rightarrow$  soft assignments; Low  $T \rightarrow$  near-hard clustering.

# Temperature Dynamics

- Attention  $a_{ik}$ :

$$a_{ik} = \frac{\exp\left(-\frac{\|x_i - \mu_k\|^2}{\tau}\right)}{\sum_j \exp\left(-\frac{\|x_i - \mu_j\|^2}{\tau}\right)}$$

- $\alpha$  (inverse temperature) controls clustering sharpness:

$$\lim_{\alpha \rightarrow 0} a_{ik} = \frac{1}{K}, \quad \lim_{\alpha \rightarrow \infty} a_{ik} = \delta_{k, \arg \min_j \|x_i - \mu_j\|^2}$$

- Smooth interpolation between soft and hard assignments.

# Temperature Dynamics

- Attention  $a_{ik}$ :

$$a_{ik} = \frac{\exp\left(-\frac{\|x_i - \mu_k\|^2}{\tau}\right)}{\sum_j \exp\left(-\frac{\|x_i - \mu_j\|^2}{\tau}\right)}$$

- $\alpha$  (inverse temperature) controls clustering sharpness:

$$\lim_{\alpha \rightarrow 0} a_{ik} = \frac{1}{K}, \quad \lim_{\alpha \rightarrow \infty} a_{ik} = \delta_{k, \arg \min_j \|x_i - \mu_j\|^2}$$

- Smooth interpolation between soft and hard assignments.

## Practical Schedule

$$\alpha_t = \alpha_0 + (\alpha_{\max} - \alpha_0) \frac{t}{T_{\text{train}}}$$

Gradually increase  $\alpha$  for stable training  $\rightarrow$  sharp clustering.

## Figure: Temperature Effect

- Low  $\alpha$ : smooth assignments; High  $\alpha$ : peaked responses
- Equivalent to annealing in Gumbel-Softmax relaxations
- Following plot of soft assignment probabilities for different  $\alpha$

# Gumbel Soft-Max distribution

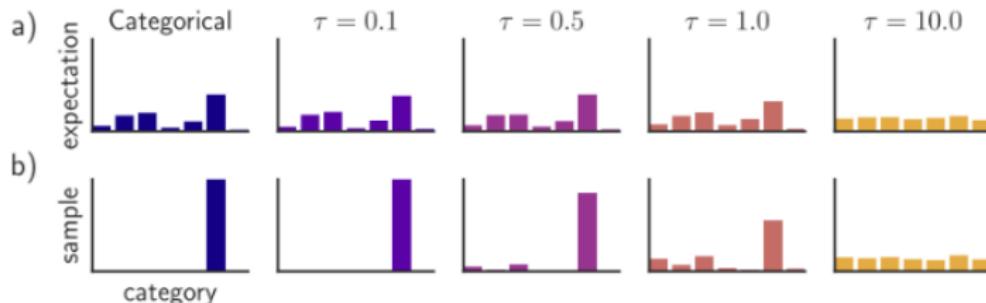
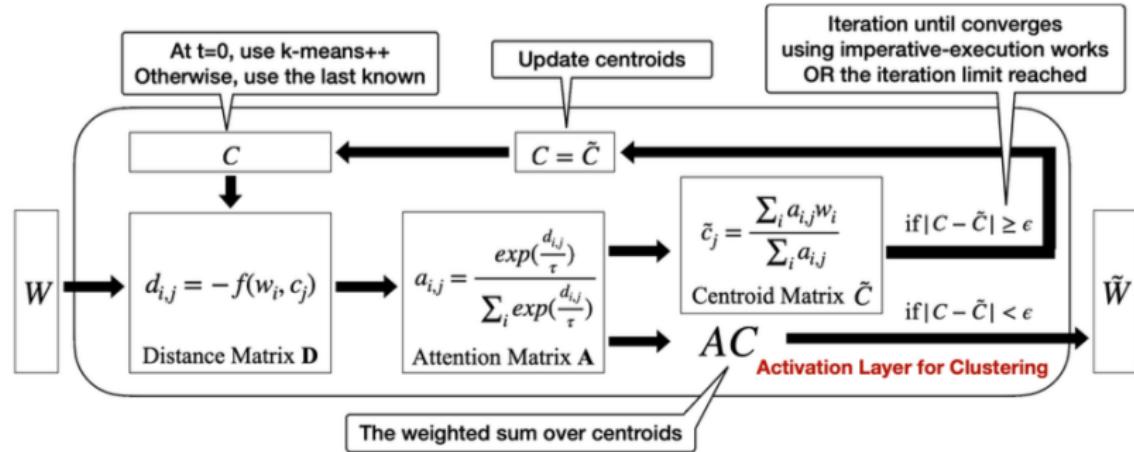


Figure 1: The Gumbel-Softmax distribution interpolates between discrete one-hot-encoded categorical distributions and continuous categorical densities. (a) For low temperatures ( $\tau = 0.1, \tau = 0.5$ ), the expected value of a Gumbel-Softmax random variable approaches the expected value of a categorical random variable with the same logits. As the temperature increases ( $\tau = 1.0, \tau = 10.0$ ), the expected value converges to a uniform distribution over the categories. (b) Samples from Gumbel-Softmax distributions are identical to samples from a categorical distribution as  $\tau \rightarrow 0$ . At higher temperatures, Gumbel-Softmax samples are no longer one-hot, and become uniform as  $\tau \rightarrow \infty$ .

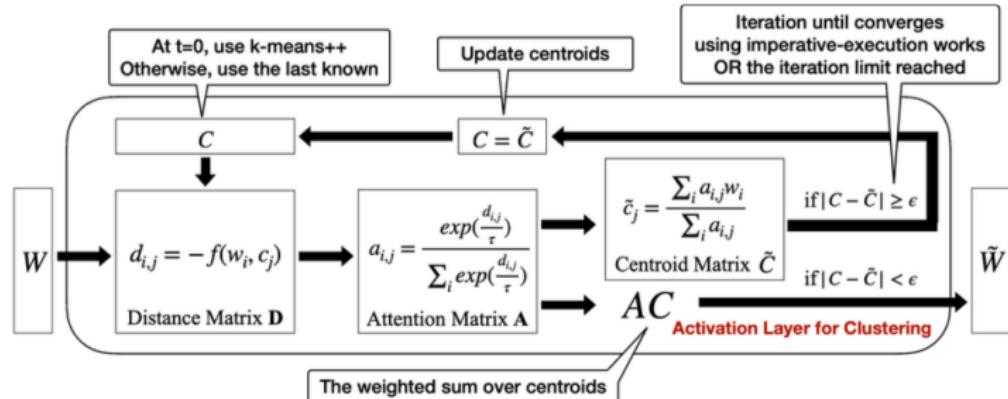
Jang et al, "Categorical Reparametrization with Gumbel-SoftMax", ICLR 2017, <https://arxiv.org/pdf/1611.01144>

# Attention-Like Soft Assignment (1)



- $C \in \mathbb{R}^K$  is the vector of cluster centers and  $W \in \mathbb{R}^N$  vector of weights
- initialization of  $C$  by randomly select K weights from  $W$  or by K-Means++
- distance matrix D computed using differentiable metric (Euclidean distance)  $f$ , as  $d_{ij} = -f(w_i, c_j)$

## Attention-Like Soft Assignment (2)



- softmax with temperature  $\tau$  applied on each row of  $D$  to get the attention matrix  $A$ ,  $a_{ij}$  from  $w_i$  to  $c_j$
  - compute centroid candidate  $\tilde{C}$  (expectation phase from EM algorithm) gathering contributions into  $\tilde{c}_j$
  - repeat process until convergence,  $|C - \tilde{C}| \leq \epsilon$
  - compute  $A \cdot C$  to get  $\tilde{W}$  for forward-propagation
  - since  $W = AC$ , the iterative process is differentiable for backprop

# Compression Ratio

- Replace each weight tensor  $W$  by its soft clustered version  $\hat{W}$ .
- Model size after clustering:

$$S_{\text{compressed}} = K \cdot b_\mu + N \cdot \log_2 K$$

where  $b_\mu$  is centroid bitwidth.

- Compression ratio:

$$R = \frac{N \cdot b_W}{S_{\text{compressed}}}$$

## Result (from Table 3)

ResNet50  $\rightarrow 29.4 \times$  compression,  $-1.6\%$  accuracy drop.

# Multi-Dimensional DKM

- $N$  elements of weights are split into  $\frac{N}{d}$  contiguous  $d$  dimensional sub-vectors
- E.g. flatten all convolutional layers into a  $(\frac{N}{d}, d)$  matrix
- Centroids  $c_j \in \mathbb{R}^d$  are computed for all subvectors  $w_i \in \mathbb{R}^d$
- $d$ -dim subvectors  $w_i$  are clustered using  $b$  bits per subvector (i.e.  $2^b$  centroids)
- $b/d$  bits for each scalar weight
- Compression Ratio for a 32-bit weight:

$$CR = \frac{32}{b/d}$$

# Multi-Dimensional DKM

- $N$  elements of weights are split into  $\frac{N}{d}$  contiguous  $d$  dimensional sub-vectors
- E.g. flatten all convolutional layers into a  $(\frac{N}{d}, d)$  matrix
- Centroids  $c_j \in \mathbb{R}^d$  are computed for all subvectors  $w_i \in \mathbb{R}^d$
- $d$ -dim subvectors  $w_i$  are clustered using  $b$  bits per subvector (i.e.  $2^b$  centroids)
- $b/d$  bits for each scalar weight
- Compression Ratio for a 32-bit weight:

$$CR = \frac{32}{b/d}$$

## Result

- Keeping same CR we can increase  $d$  while adjusting also  $d$  increase the entropy of weights' distribution
- Small vectors preserves local structure → LUT represents richer patterns

# Experimental Results

Model	Compression	Top-1	Drop
ResNet50 (ImageNet)	29.4×	74.5%	-1.6%
MobileNetV1	22.4×	63.9%	-7.0%
DistilBERT	11.8×	83.5%	-1.1%

## Observation

- DKM vs. Quantization, Pruning, and Post-hoc Clustering
- Jointly learned centroids outperform discrete quantization
- DKM achieves high compression with minimal loss — outperforming existing clustering methods

# Comparison with SotA

	Base (32 bit)	ResNet18	ResNet50	MobileNet-v1	MobileNet-v2
3 bit	PROFIT			69.6	69.6
	EWGS	<b>70.5</b>	<b>76.3</b>	64.4	64.5
	PROFIT+EWGS			68.6	69.5
	DKM	69.9	76.2	<b>69.9</b>	<b>70.3</b>
2 bit	PROFIT			63.4	61.9
	EWGS	<b>69.3</b>	<b>75.8</b>	52.0	49.1
	DKM	68.9	75.3	<b>66.4</b>	<b>66.2</b>
1 bit	PROFIT			nc <sup>o</sup>	nc
	EWGS	66.6	73.8	8.5	23.0
	DKM 1/1	65.0	72.1	5.9	50.8
	DKM 4/4 <sup>§</sup>	67.0	<b>73.8</b>	60.6	55.0
	DKM 8/8	<b>67.8</b>	oom <sup>□</sup>	<b>64.3</b>	<b>62.4</b>
$\frac{1}{2}$ bit	DKM 4/8	62.1	70.6	46.5	34.0
	DKM 8/16	<b>65.5</b>	<b>72.1</b>	<b>59.8</b>	<b>58.3</b>

<sup>o</sup> not converging; <sup>□</sup> out of memory ; <sup>§</sup> clustering with 4 bits and 4 dimensions

**Table 1:** When compared with the latest weight quantization algorithms, DKM-based algorithm shows superior Top-1 accuracy when the network is hard to optimize (i.e., MobileNet-v1/v2) or when a low precision is required (1 bit). Further, with multi-dimensional DKM (see Section 3.3), DKM delivers 64.3 % Top-1 accuracy for MobileNet-v1 with the 8/8 configuration which is equivalent to 1 bit-per-weight.

# Improved BERT performance

		ALBERT	DistilBERT	BERT-tiny	MobileBERT
	Base (32 bit)	90.6	88.2	78.9	89.6
3 bit	EWGS	83.3	87.6	78.3	87.8
	DKM	<b>85.1</b>	<b>88.2</b>	<b>80.0</b>	<b>89.0</b>
2 bit	EWGS	79.6	85.4	77.9	81.6
	DKM	<b>81.7</b>	<b>87.4</b>	<b>80.0</b>	<b>83.7</b>
1 bit	EWGS	62.0	60.9	74.5	60.2
	DKM	79.0	82.8	<b>77.4</b>	69.8
	DKM 4/4 <sup>§</sup>	<b>80.0</b>	<b>84.0</b>	77.2	<b>78.3</b>

<sup>§</sup> clustering with 4 bits and 4 dimensions

# Key Takeaways

- DKM introduces a differentiable relaxation of K-Means via temperature-controlled soft assignments.
- Temperature ( $T = 1/\alpha$ ) is crucial:
  - High  $T \rightarrow$  stable gradients, fuzzy clustering
  - Low  $T \rightarrow$  crisp clustering, high compression
- Proper annealing schedule ensures smooth convergence.
- Memory-heavy to train:

$$O(N + Kd + NK)$$

Last term is the (implicit) attention map, computed during backpropagation

# References

-  Y. Cho et al., "*DKM: Differentiable K-Means Clustering Layer for Neural Network Compression*," arXiv:2108.12659, 2021
-  S. Lloyd, "*Least squares quantization in PCM*," IEEE Trans. Inf. Theory, 1982
-  E. Jang et al., "*Categorical reparameterization with Gumbel–Softmax*," ICLR 2017
-  S. Jaffe et al. "*IDKM: Memory Efficient Neural Network Quantization via Implicit, Differentiable k-Means*" ICLR 2023