


Elegantly Configured ML Training Frameworks with Lightning & Hydra

Ceanga (Rădulescu) Teodora - 20.11.2025

PyTorch Lightning

<https://lightning.ai/docs/pytorch/stable/>

- A Deep learning framework built on top of PyTorch for researchers & engineers
- Faster iterations than vanilla PyTorch
- Key advantages
 - Organization (standardization , decoupling)
 - Scalability (multi GPU, TPU)
 - Reproducibility



Hydra

<https://hydra.cc/>

- A framework for elegantly configuring complex applications
- Key advantages:
 - No boilerplate
 - Powerful configuration
 - Pluggable architecture








Lightning + Hydra

<https://github.com/ashleve/lightning-hydra-template>

- **Rapid Experimentation**: thanks to hydra command line superpowers
- **Minimal Boilerplate**: thanks to automating pipelines with config instantiation
- **Main Configs**: allow you to specify default training configuration
- **Experiment Configs**: allow you to override chosen hyperparameters and version control experiments
- **Workflow**: comes down to 4 simple steps
- **Experiment Tracking**: Tensorboard, W&B, Neptune, Comet, MLFlow and CSVLogger
- **Logs**: all logs (checkpoints, configs, etc.) are stored in a dynamically generated folder structure
- **Hyperparameter Search**: simple search is effortless with Hydra plugins like Optuna Sweeper
- **Tests**: generic, easy-to-adapt smoke tests for speeding up the development
- **Continuous Integration**: automatically test and lint your repo with Github Actions
- **Best Practices**: a couple of recommended tools, practices and standards

PyTorch → Lightning



```
def train_loop(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    # Set the model to training mode - important for batch normalization and dropout layers
    # Unnecessary in this situation but added for best practices
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        # Compute prediction and loss
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()
        torch.optim.SGD.step()

        if batch % 100 == 0:
            loss, current = loss.item(), batch * batch_size + len(X)
            print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

def test_loop(dataloader, model, loss_fn):
    # Set the model to evaluation mode - important for batch normalization and dropout layers
    # Unnecessary in this situation but added for best practices
    model.eval()
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    test_loss, correct = 0, 0

    # Evaluating the model with torch.no_grad() ensures that no gradients are computed during test mode
    # also serves to reduce unnecessary gradient computations and memory usage for tensors with requires_grad=True
    with torch.no_grad():
        for X, y in dataloader:
            pred = model(X)
            test_loss += loss_fn(pred, y).item()
            correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    test_loss /= num_batches
    correct /= size
    print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>8f} \n")

    loss_fn = nn.CrossEntropyLoss()
    optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)

    epochs = 10
    for t in range(epochs):
        print(f"Epoch {t+1}\n-----")
        train_loop(train_dataloader, model, loss_fn, optimizer)
        test_loop(test_dataloader, model, loss_fn)
    print("Done!")
```

Lightning DataModule

<https://lightning.ai/docs/pytorch/latest/data/datamodule.html>

A LightningDataModule implements 7 key methods:

- **def prepare_data(self):**
 - Things to do on 1 GPU/TPU (not on every GPU/TPU in DDP).
 - Download data, pre-process, split, save to disk, etc...
- **def setup(self, stage):**
 - Things to do on every process in DDP.
 - Load data, set variables, etc...
- **def train_dataloader(self):**
- **def val_dataloader(self):**
- **def test_dataloader(self):**
- **def predict_dataloader(self):**
- **def teardown(self, stage):**
 - Called on every process in DDP.
 - Clean up after fit or test

Lightning Module


https://lightning.ai/docs/pytorch/latest/common/lightning_module.html

A `LightningModule` implements 8 key methods:

- **def __init__(self):**
 - Define initialization code here.
- **def setup(self, stage):**
 - Things to setup before each stage, 'fit', 'validate', 'test', 'predict'.
 - This hook is called on every process when using DDP.
- **def training_step(self, batch, batch_idx):**
 - The complete training step.
- **def validation_step(self, batch, batch_idx):**
 - The complete validation step.
- **def test_step(self, batch, batch_idx):**
 - The complete test step.
- **def predict_step(self, batch, batch_idx):**
 - The complete predict step.
- **def configure_optimizers(self):**
 - Define and configure optimizers and LR schedulers.

Lightning Trainer

<https://lightning.ai/docs/pytorch/latest/common/trainer.html>



- Under the hood the Trainer handles:
 - Enabling / disabling gradients
 - Iterating through the data loaders
 - Putting data & computations on correct devices
 - Backpropagation & parameter updates
 - Optimizer and LR Scheduler calls

Lightning Callbacks

<https://lightning.ai/docs/pytorch/latest/extensions/callbacks.html>

- The training loop automated by the Trainer can be customized using Callbacks.
- Callbacks are arbitrary self-contained programs that can be added to the training (for customizing things like checkpointing, model summary, progress bars, etc) using **hooks**
- Common hooks:
 - `on_load_checkpoint`
 - `on_validation_epoch_start`, `on_training_epoch_end`
 - `on_before_backward`, `on_after_backward`

Hydra - getting started

```
pip install hydra-core --upgrade
```

```
src > ⏎ train copy.py > ...
1 import hydra
2 from lightning import LightningDataModule, LightningModule, Trainer
3 from omegaconf import DictConfig
4
5
6
7 def train(cfg: DictConfig):
8
9     datamodule: LightningDataModule = hydra.utils.instantiate(cfg.data)
10
11    model: LightningModule = hydra.utils.instantiate(cfg.model)
12
13    trainer: Trainer = hydra.utils.instantiate(cfg.trainer)
14
15    if cfg.get("train"):
16        |   trainer.fit(model=model, datamodule=datamodule, ckpt_path=cfg.get("ckpt_path"))
17
18    if cfg.get("test"):
19        ckpt_path = trainer.checkpoint_callback.best_model_path
20        trainer.test(model=model, datamodule=datamodule, ckpt_path=ckpt_path)
21
22
23 @hydra.main(version_base="1.3", config_path="../configs", config_name="train.yaml")
24 def main(cfg: DictConfig):
25     """Main entry point for training.
26
27     :param cfg: DictConfig configuration composed by Hydra.
28     """
29     train(cfg)
30
31
32 if __name__ == "__main__":
33     main()
```

Basic folder structure:

- ◆ my_project/
 - configs/
 - train.yaml
 - src/
 - data/
 - models/
 - train.py

Hydra - example config

```
configs > ! train copy.yaml
1 # specify here default configuration
2 # order of defaults determines the order in which configs override each other
3 defaults:
4   - _self_
5   - data: mnist
6   - model: mnist
7   - trainer: default
8
9 # experiment configs allow for version control of specific hyperparameters
10 # e.g. best hyperparameters for given model and datamodule
11 - experiment: null
12
13 # set False to skip model training
14 train: True
15
16 # evaluate on test set, using best model weights achieved during training
17 # lightning chooses best weights based on the metric specified in checkpoint callback
18 test: True
19
20 # simply provide checkpoint path to resume training
21 ckpt_path: null
22
```

```
configs > trainer > ! default.yaml
1 _target_: lightning.pytorch.trainer.Trainer
2
3 default_root_dir: ${paths.output_dir}
4
5 min_epochs: 1 # prevents early stopping
6 max_epochs: 10
7
8 accelerator: cpu
9 devices: 1
10
11 # mixed precision for extra speed-up
12 # precision: 16
13
14 # perform a validation loop every N training epochs
15 check_val_every_n_epoch: 1
16
17 # set True to to ensure deterministic results
18 # makes training slower but gives more reproducibility than just setting seeds
19 deterministic: False
20
```

```
configs > data > ! mnist.yaml
1 _target_: src.data.mnist_datamodule.MNISTDataModule
2 data_dir: ${paths.data_dir}
3 batch_size: 128 # Needs to be divisible by the number of devices (e.g., if in a distributed setup)
4 train_val_test_split: [55_000, 5_000, 10_000]
5 num_workers: 0
6 pin_memory: False
7
```

Value interpolation

```
configs > model > ! mnist.yaml
1 _target_: src.models.mnist_module.MNISTLitModule
2
3 optimizer:
4   _target_: torch.optim.Adam
5   _partial_: true
6   lr: 0.001
7
8 scheduler:
9   _target_: torch.optim.lr_scheduler.ReduceLROnPlateau
10  _partial_: true
11  mode: min
12  factor: 0.1
13  patience: 10
14
15 net:
16   _target_: src.models.components.simple_dense_net.SimpleDenseNet
17   input_size: 784
18   lin1_size: 64
19   lin2_size: 128
20   lin3_size: 64
21   output_size: 10
22
23 # compile model for faster training with pytorch 2.0
24 compile: false
25
```

Hydra - advanced features

- **Value interpolation**
 - https://hydra.cc/docs/upgrades/1.0_to_1.1/defaults_list_interpolation/
- **Custom resolvers**
 - https://omegaconf.readthedocs.io/en/2.1_branch/custom_resolvers.html
- **Partial instantiation**
 - https://hydra.cc/docs/advanced/instantiate_objects/overview/

Lightning-hydra-template

Code walkthrough

<https://github.com/ashleve/lightning-hydra-template>

Cool Tricks

```
! mnist.yaml •  
configs > model > ! mnist.yaml  
1   _target_: src.models.mnist_module.MNISTLitModule  
2  
3   optimizer:  
4     _target_: torch.optim.Adam  
5     _partial_: true  
6     lr: 0.001  
7  
8   scheduler:  
9     _target_: torch.optim.lr_scheduler.ReduceLROnPlateau  
10    _partial_: true  
11    mode: min  
12    factor: 0.1  
13    patience: 10
```



Override any config



Add new parameters with + sign

```
python train.py +model.optimizer.weight_decay=0.1
```



Accelerators and distributed training

```
# train on CPU
python train.py trainer.accelerator=cpu

# train on 1 GPU
python train.py trainer.accelerator=gpu

# train on TPU
python train.py +trainer.tpu_cores=8

# train with DDP (Distributed Data Parallel) (4 GPUs)
python train.py trainer.accelerator=gpu trainer.strategy=ddp trainer.devices=4

# train with DDP (Distributed Data Parallel) (8 GPUs, 2 nodes)
python train.py trainer.accelerator=gpu trainer.strategy=ddp trainer.devices=4 trainer.num_nodes=2

# simulate DDP on CPU processes
python train.py trainer.accelerator=cpu trainer.strategy=ddp_spawn trainer.devices=2

# accelerate training on mac
python train.py trainer.accelerator=mps
```

Other tricks

Training / Optimization

- `max_epochs`
- `min_epochs`
- `max_steps`
- `limit_train_batches`
- `limit_val_batches`
- `limit_test_batches`
- `limit_predict_batches`
- `gradient_clip_val`
- `gradient_clip_algorithm`
- `accumulate_grad_batches`
- `deterministic`
- `benchmark`

Devices & Distributed Training

- `devices` (e.g., "auto", 1, 4, [0, 1])
- `accelerator` ("cpu", "gpu", "cuda", "mps", "tpu", "auto")
- `strategy` ("ddp", "deepspeed", "fsdp", "dp", "auto")
- `num_nodes`
- `precision` (16, 32, bf16, "16-mixed", etc.)
- `sync_batchnorm`
- .

Logging / Checkpointing

- `logger`
- `enable_checkpointing`
- `default_root_dir`
- `log_every_n_steps`
- `enable_progress_bar`
- `enable_model_summary`
- `callbacks`

Validation & Early Stopping

- `val_check_interval`
- `check_val_every_n_epoch`
- `enable_validation`
- `reload_dataloaders_every_n_epochs`